

# ECE 414/517

# Reinforcement Learning

LECTURE 8: DYNAMIC PROGRAMMING

SEP. 20 2022

(ADAPTED FROM SLIDES BY DR. ITAMAR AREL)

# Dynamic Programming

---

Dynamic programming is the collection of algorithms that can be used to compute optimal policies given perfect model of the environment.

1. DP constitutes a theoretically optimal methodology
2. In reality it is often limited since DP is computationally expensive.
3. As related to other methods:
  - Hope to do “just as well” as DP
  - Hopefully can do it with less computation
  - Hopefully can do it with less memory
  - Hopefully without model
4. Key Idea (In general): Use value function to derive optimal policy
5. Need to find a way to calculate optimal state value function

# Policy Iteration

1. Technique for obtaining the optimal policy

2. Comprises of two complementing steps:


- Policy Evaluation: updating the value function in view of current policy

- $V^\pi(s) = \sum_a \pi(a|s) \sum_{s',r} p(s',r|s,a)[r + \gamma V^\pi(s')]$




- $V_{k+1}^\pi(s) = \sum_a \pi(a|s) \sum_{s',r} p(s',r|s,a)[r + \gamma V_k^\pi(s')]$
- Policy Improvement: updating the policy given the current value function (which can be sub-optimal)
  - $\pi'(s) = \arg \max_a \sum_{s',r} p(s',r|s,a)[r + \gamma V^\pi(s')]$
- It can be shown that  $\pi'$  is at least as good as  $\pi$ , and if they are equal they are both the optimal policy.

$$\pi_0 \xrightarrow{E} V_{\pi_0} \xrightarrow{I} \pi_1 \xrightarrow{E} V_{\pi_1} \xrightarrow{I} \pi_2 \xrightarrow{E} \dots \xrightarrow{I} \pi_* \xrightarrow{E} V_{\pi_*}$$



# Example: Grid World – Policy Evaluation



$$\pi(s) = \{0.25, 0.25, 0.25, 0.25\} \forall s \in \mathcal{S}$$

$$\gamma = 0.9$$

$$V_{k+1}^{\pi}(s) = \sum_a \pi(a|s) \sum_{s', r} p(s', r|s, a) [r + \gamma V_k^{\pi}(s')]$$

3.31	8.79	4.43	5.32	1.49
1.52	2.99	2.25	1.91	0.55
0.05	0.74	0.67	0.36	-0.40
-0.97	-0.44	-0.35	-0.59	-1.18
-1.86	-1.34	-1.23	-1.42	-1.97

$V_{\pi_0}$

0	0	0	0	0
0	0	0	0	0
0	0	0	0	0
0	0	0	0	0
0	0	0	0	0

$V_0$

-0.50	10.0	-0.25	5.00	-0.5
-0.25	0.00	0.00	0.00	-0.25
-0.25	0.00	0.00	0.00	-0.25
-0.25	0.00	0.00	0.00	-0.25
-0.50	-0.25	-0.25	-0.25	-0.5

$V_1$

1.47	9.78	3.07	5.00	0.34
-0.48	2.19	-0.06	1.07	-0.48
-0.42	-0.06	0.00	-0.06	-0.42
-0.48	-0.11	-0.06	-0.11	-0.48
-0.84	-0.48	-0.42	-0.48	-0.84


$V_2$

2.25	9.57	3.75	4.95	0.67
0.37	2.07	1.42	0.99	-0.13
-0.57	0.37	-0.05	0.12	-0.57
-0.66	-0.24	-0.14	-0.24	-0.66
-1.09	-0.66	-0.57	-0.66	-1.09

$V_3$

$$\pi_0 \xrightarrow{E} V_{\pi_0} \xrightarrow{I} \pi_1 \xrightarrow{E} V_{\pi_1} \xrightarrow{I} \pi_2 \xrightarrow{E} \dots \xrightarrow{I} \pi_* \xrightarrow{E} V_{\pi_*}$$

# Example: Grid World – Policy Improvement



$$\pi_0(s) = \{0.25, 0.25, 0.25, 0.25\} \forall s \in S$$

$$\gamma = 0.9$$

$$\pi'(s) = \arg \max_a \sum_{s', r} p(s', r | s, a) [r + \gamma V^\pi(s')]$$

3.31	8.79	4.43	5.32	1.49
1.52	2.99	2.25	1.91	0.55
0.05	0.74	0.67	0.36	-0.40
-0.97	-0.44	-0.35	-0.59	-1.18
-1.86	-1.34	-1.23	-1.42	-1.97

$V_{\pi_0}$

→	↓	←	↓	←
↑	↑	↑	↑	←
↑	↑	↑	↑	↑
↑	↑	↑	↑	↑
↑	↑	↑	↑	↑

$\pi_1$

$$\pi(0,0) = \arg \max \begin{bmatrix} -1 + 0.9 \cdot 3.31 \\ -1 + 0.9 \cdot 3.31 \\ 0 + 0.9 \cdot 8.79 \\ 0 + 0.9 \cdot 1.52 \end{bmatrix} \begin{matrix} W \\ N \\ E \\ S \end{matrix} = E$$

$$= \{0, 0, 1, 0\}$$

$$\pi_0 \xrightarrow{E} V_{\pi_0} \xrightarrow{I} \pi_1 \xrightarrow{E} V_{\pi_1} \xrightarrow{I} \pi_2 \xrightarrow{E} \dots \xrightarrow{I} \pi_* \xrightarrow{E} V_{\pi_*}$$

# Policy Improvement

---

$$\begin{aligned} v_{\pi}(s) &\leq q_{\pi}(s, \pi'(s)) \\ &= \mathbb{E}[R_{t+1} + \gamma v_{\pi}(S_{t+1}) \mid S_t = s, A_t = \pi'(s)] \\ &= \mathbb{E}_{\pi'}[R_{t+1} + \gamma v_{\pi}(S_{t+1}) \mid S_t = s] \\ &\leq \mathbb{E}_{\pi'}[R_{t+1} + \gamma q_{\pi}(S_{t+1}, \pi'(S_{t+1})) \mid S_t = s] \\ &= \mathbb{E}_{\pi'}[R_{t+1} + \gamma \mathbb{E}_{\pi'}[R_{t+2} + \gamma v_{\pi}(S_{t+2}) \mid S_{t+1}] \mid S_t = s] \\ &= \mathbb{E}_{\pi'}[R_{t+1} + \gamma R_{t+2} + \gamma^2 v_{\pi}(S_{t+2}) \mid S_t = s] \\ &\leq \mathbb{E}_{\pi'}[R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \gamma^3 v_{\pi}(S_{t+3}) \mid S_t = s] \\ &\vdots \\ &\leq \mathbb{E}_{\pi'}[R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \gamma^3 R_{t+4} + \cdots \mid S_t = s] \\ &= v_{\pi'}(s). \end{aligned}$$

# Policy Iteration

Policy Iteration (using iterative policy evaluation) for estimating  $\pi \approx \pi_*$

1. Initialization

$V(s) \in \mathbb{R}$  and  $\pi(s) \in \mathcal{A}(s)$  arbitrarily for all  $s \in \mathcal{S}$

2. Policy Evaluation

Loop:

$\Delta \leftarrow 0$

Loop for each  $s \in \mathcal{S}$ :

$v \leftarrow V(s)$

$V(s) \leftarrow \sum_{s',r} p(s', r | s, \pi(s)) [r + \gamma V(s')]$

$\Delta \leftarrow \max(\Delta, |v - V(s)|)$

until  $\Delta < \theta$  (a small positive number determining the accuracy of estimation)

3. Policy Improvement

*policy-stable*  $\leftarrow$  true

For each  $s \in \mathcal{S}$ :


*old-action*  $\leftarrow \pi(s)$

$\pi(s) \leftarrow \operatorname{argmax}_a \sum_{s',r} p(s', r | s, a) [r + \gamma V(s')]$

If *old-action*  $\neq \pi(s)$ , then *policy-stable*  $\leftarrow$  false

If *policy-stable*, then stop and return  $V \approx v_*$  and  $\pi \approx \pi_*$ ; else go to 2

# Example: Grid World – Policy Evaluation



$$\pi_0(s) = \{0.25, 0.25, 0.25, 0.25\} \forall s \in S$$

$$\gamma = 0.9$$

$$V_{k+1}^\pi(s) = \sum_a \pi(a|s) \sum_{s', r} p(s', r|s, a) [r + \gamma V_k^\pi(s')]$$

3.31	8.79	4.43	5.32	1.49
1.52	2.99	2.25	1.91	0.55
0.05	0.74	0.67	0.36	-0.40
-0.97	-0.44	-0.35	-0.59	-1.18
-1.86	-1.34	-1.23	-1.42	-1.97

$V_{\pi_0}$

→	↓	←	↓	←
↑	↑	↑	↑	←
↑	↑	↑	↑	↑
↑	↑	↑	↑	↑
↑	↑	↑	↑	↑

$\pi_1$


21.98	24.42	21.98	19.42	17.48
19.78	21.98	19.78	17.80	16.02
17.80	19.78	17.80	16.02	14.42
16.02	17.80	16.02	14.42	12.98
14.42	16.02	14.42	12.98	11.68

$V_{\pi_1}$

$$\pi_0 \xrightarrow{E} V_{\pi_0} \xrightarrow{I} \pi_1 \xrightarrow{E} V_{\pi_1} \xrightarrow{I} \pi_2 \xrightarrow{E} \dots \xrightarrow{I} \pi_* \xrightarrow{E} V_{\pi_*}$$



# Example: Grid World – Policy Improvement



$$\pi_0(s) = \{0.25, 0.25, 0.25, 0.25\} \forall s \in S$$

$$\gamma = 0.9$$

$$\pi'(s) = \arg \max_a \sum_{s', r} p(s', r | s, a) [r + \gamma V^\pi(s')]$$

3.31	8.79	4.43	5.32	1.49
1.52	2.99	2.25	1.91	0.55
0.05	0.74	0.67	0.36	-0.40
-0.97	-0.44	-0.35	-0.59	-1.18
-1.86	-1.34	-1.23	-1.42	-1.97

$V_{\pi_0}$

→	↓	←	↓	←
↑	↑	↑	↑	←
↑	↑	↑	↑	↑
↑	↑	↑	↑	↑
↑	↑	↑	↑	↑

$\pi_1$

21.98	24.42	21.98	19.42	17.48
19.78	21.98	19.78	17.80	16.02
17.80	19.78	17.80	16.02	14.42
16.02	17.80	16.02	14.42	12.98
14.42	16.02	14.42	12.98	11.68

$V_{\pi_1}$

→	↓	←	↓	←
↑	↑	↑	↑	←
↑	↑	↑	↑	↑
↑	↑	↑	↑	↑
↑	↑	↑	↑	↑

$\pi_2$

$$\pi_0 \xrightarrow{E} V_{\pi_0} \xrightarrow{I} \pi_1 \xrightarrow{E} V_{\pi_1} \xrightarrow{I} \pi_2 \xrightarrow{E} \dots \xrightarrow{I} \pi_* \xrightarrow{E} V_{\pi_*}$$

# Value Iteration

---

$$V_{k+1}^{\pi}(s) = \sum_a \pi(a|s) \sum_{s',r} p(s',r|s,a)[r + \alpha V_k^{\pi}(s')]$$

$$V_0 \rightarrow V_1 \rightarrow V_2 \rightarrow \dots \rightarrow V_k \rightarrow V_{k+1} \rightarrow \dots \rightarrow V^{\pi}$$

$$\pi'(s) = \arg \max_a \sum_{s',r} p(s',r|s,a)[r + \alpha V^{\pi}(s')]$$

$$\pi_0 \xrightarrow{E} V_{\pi_0} \xrightarrow{I} \pi_1 \xrightarrow{E} V_{\pi_1} \xrightarrow{I} \pi_2 \xrightarrow{E} \dots \xrightarrow{I} \pi_* \xrightarrow{E} V_{\pi_*}$$



$$V_{k+1}(s) = \max_a \sum_{s',r} p(s',r|s,a)[r + \alpha V_k(s')]$$

Solving the Bellman  
Optimality Equation  
Iteratively

# Value Iteration

## Value Iteration, for estimating $\pi \approx \pi_*$

Algorithm parameter: a small threshold  $\theta > 0$  determining accuracy of estimation

Initialize  $V(s)$ , for all  $s \in \mathcal{S}^+$ , arbitrarily except that  $V(\text{terminal}) = 0$

Loop:


```
|  $\Delta \leftarrow 0$   
| Loop for each  $s \in \mathcal{S}$ :  
|    $v \leftarrow V(s)$   
|    $V(s) \leftarrow \max_a \sum_{s',r} p(s', r | s, a) [r + \gamma V(s')]$   
|    $\Delta \leftarrow \max(\Delta, |v - V(s)|)$ 
```

until  $\Delta < \theta$

Output a deterministic policy,  $\pi \approx \pi_*$ , such that

$$\pi(s) = \operatorname{argmax}_a \sum_{s',r} p(s', r | s, a) [r + \gamma V(s')]$$

# Example: Grid World



$$\pi(s) = \{0.25, 0.25, 0.25, 0.25\} \forall s \in S$$

$$\gamma = 0.9$$

$$V_{k+1}^{\pi}(s) = \sum_a \pi(a|s) \sum_{s', r} p(s', r|s, a) [r + \alpha V_k^{\pi}(s')]$$

0	0	0	0	0
0	0	0	0	0
0	0	0	0	0
0	0	0	0	0
0	0	0	0	0

$V_0$

$$V_{k+1}(s) = \max_a \sum_{s', r} p(s', r|s, a) [r + \alpha V_k(s')]$$

$V_1$

-0.50	10.0	2.00	5.00	0.63
-0.36	2.17	0.94	1.34	0.19
-0.33	0.41	0.30	0.37	-0.12
-0.32	0.02	0.07	0.10	-0.26
-0.57	-0.37	-0.32	-0.30	-0.62

$V_1$

0.00	10.00	9.00	5.00	4.50
0.00	9.00	8.10	7.29	6.56
0.00	8.10	7.29	6.56	5.90
0.00	7.29	6.56	5.90	5.31
0.00	6.56	5.90	5.31	4.78

$$V_1(0,0) = \max_a \begin{bmatrix} -1 + 0.9 \cdot 0 \\ -1 + 0.9 \cdot 0 \\ 0 + 0.9 \cdot 0 \\ 0 + 0.9 \cdot 0 \end{bmatrix} \begin{matrix} W \\ N \\ E \\ S \end{matrix} = 0$$

$$V_1(0,1) = 10$$


$$v_2(0,2) = \begin{bmatrix} 0 + 0.9 \cdot 10 \\ -1 + 0.9 \cdot 0 \\ 0 + 0.9 \cdot 0 \\ 0 + 0.9 \cdot 0 \end{bmatrix} = 9$$

# Recycling Robot

h: high  
l: low  
s: search  
w: wait  
re: recharge

$$V^*(h) = \max \left\{ \begin{array}{l} R^{search} + \alpha\gamma V^*(h) + (1-\alpha)\gamma V^*(l) \\ R^{wait} + \gamma V^*(h) \end{array} \right\}$$

$$V^*(l) = \max \left\{ \begin{array}{l} \beta[R^{search} + \gamma V^*(l)] + (1-\beta)[-3 + \gamma V^*(h)] \\ R^{wait} + \gamma V^*(l) \\ \gamma V^*(h) \end{array} \right\}$$



# Gambler's Problem

1. Make bets on the outcomes of a sequence of coin flips
2. If the coin comes up heads, win as many dollars as she staked on that flip.
3. If tails, she loses her stake.
4. The game ends by reaching the goal of \$100, or losing by running out of money.
5. On each flip the gambler must decide how much of his capital to stake.
6. This problem could be formulated as an undiscounted, finite (non-deterministic MDP).
  - States?  $\{0, \dots, 100\}$  (how much \$)
  - Actions?  $A(s) = [1, \dots, s]$  (how much \$ to bet)
  - Rewards?  $+1$  win  $\&$   $0$  lose or other action
7. If probability of heads is known can be solved using dynamic programming

$$V_{k+1}(s) = \max_a \sum_{s', r} p(s', r | s, a) [r + \alpha V_k(s')]$$



$$P_h = 0.4$$

$$V_1(0) = 0 \quad (\text{bc terminal state})$$

$$\vdots V_1(1) = [0.6 \cdot (0+0) + 0.4 \cdot (0+0)] = 0$$

$$\vdots V_1(2) = \max_a \begin{bmatrix} 0 \\ 0 \end{bmatrix}$$


$$\vdots V_1(50) = \begin{bmatrix} \vdots \\ 0.6 \cdot (0+0) + 0.4 \cdot (1+0) = 0.4 \end{bmatrix} = 0.4 \quad (50)$$

$$V_1(51) = \begin{bmatrix} 0.6(0+0.4) + 0.4(0.0) \\ 0.6 \cdot (0+0) + 0.4(1+0) \end{bmatrix} \begin{matrix} (1) \\ (49) \end{matrix} = 0.4$$

$$V_1(100) = 0$$



# Gambler's Problem



$$p_h = 0.4$$

# Efficiency of Dynamic Programming

---

1. DP may not be practical for large problems (even when environment is known).
2. Still a lot better than exhaustive search.
3. If  $|A| = M$  and  $|S| = N$  then  $DP$  is guaranteed to find the optimal policy in less than some polynomial function of  $M$  and  $N$
4. Exhaustive Search:
  1.  $M^N$
5. Usually impractical because  $|S|$  grows exponentially with the number of state variables (curse of dimensionality)
6. These days DP methods can solve (using PC) MDP's with millions of states.